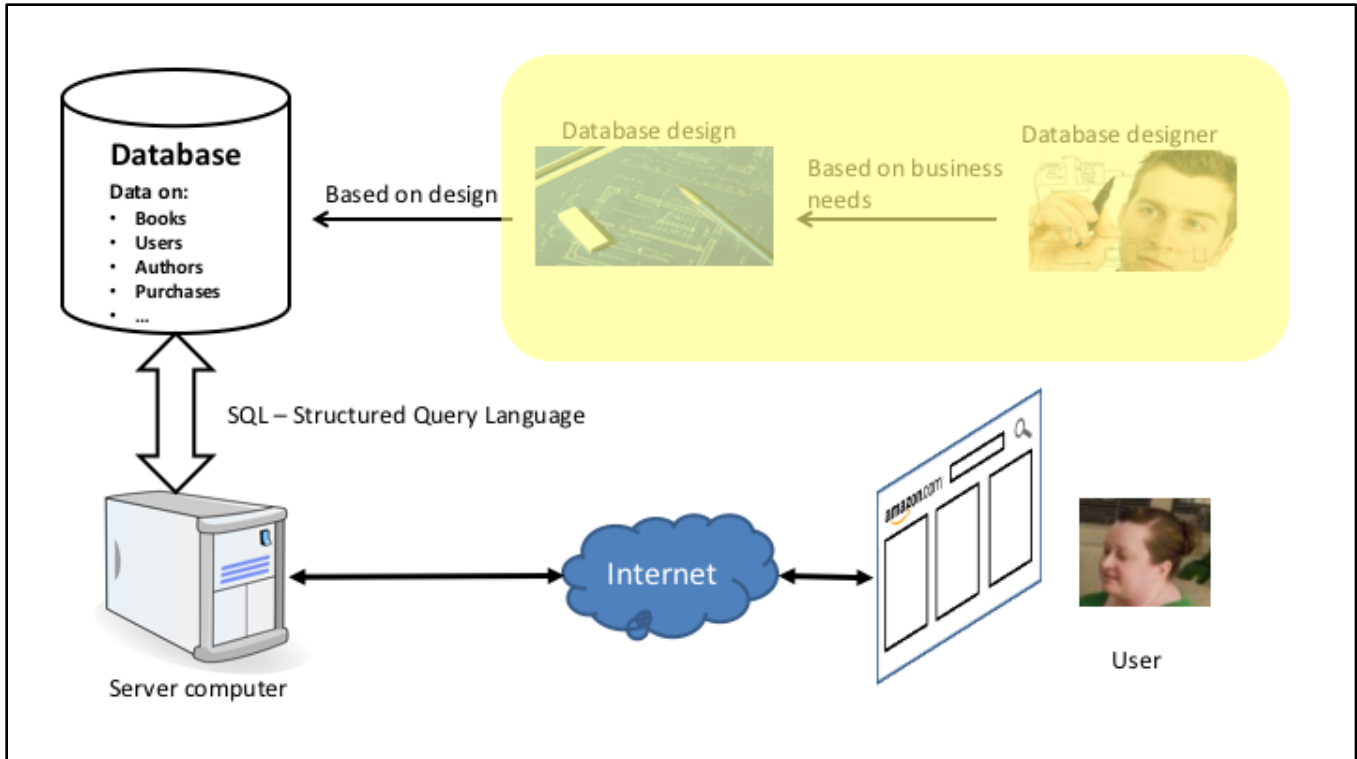


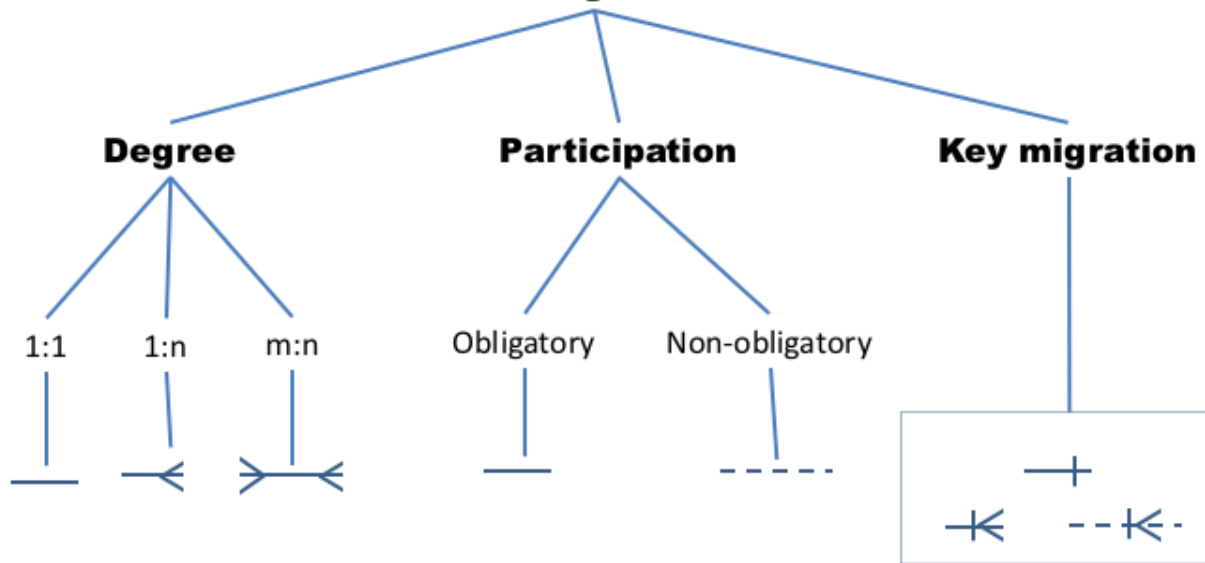
## **Database Design: Part 2**



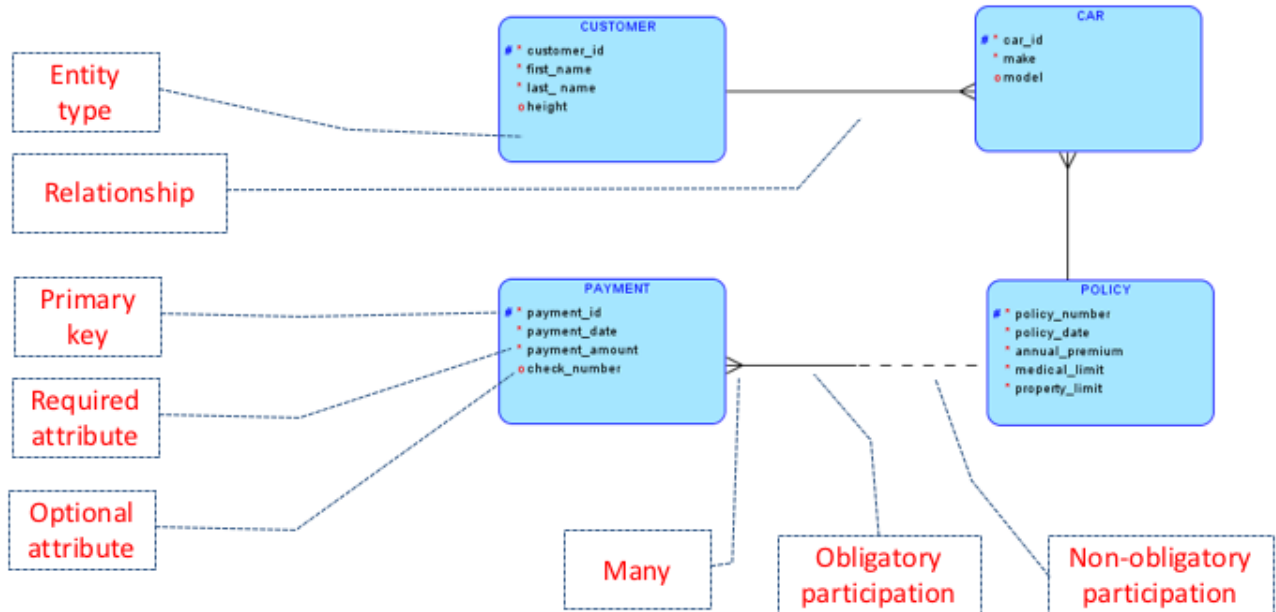
We have already looked at how SQL can help us to retrieve almost any information we need from a relational database.

We have started looking at how we can arrive at a database design in the form of an Entity Relationship Diagram for a given business situation.

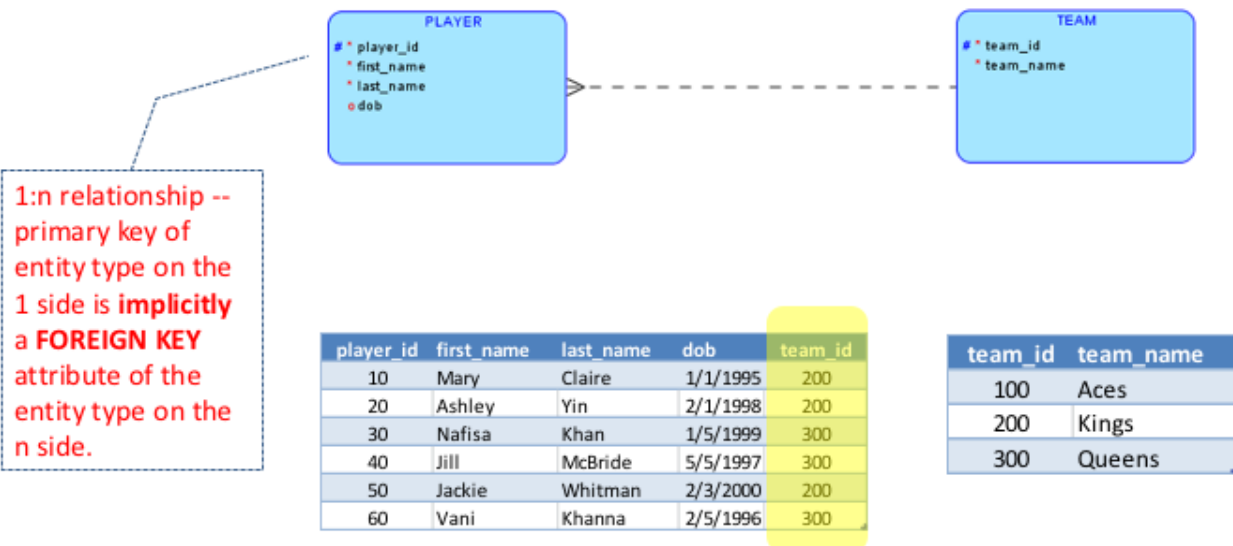
# Cardinality Notation



## Recap 1 -- Basics



## Recap 2 -- Implicit Foreign Key in 1:n Relationship



Whenever we have a 1:n relationship, we can see that representing the relationship in tables requires adding the primary key of the entity type on the 1 side as a foreign key to the entity type on the n side.

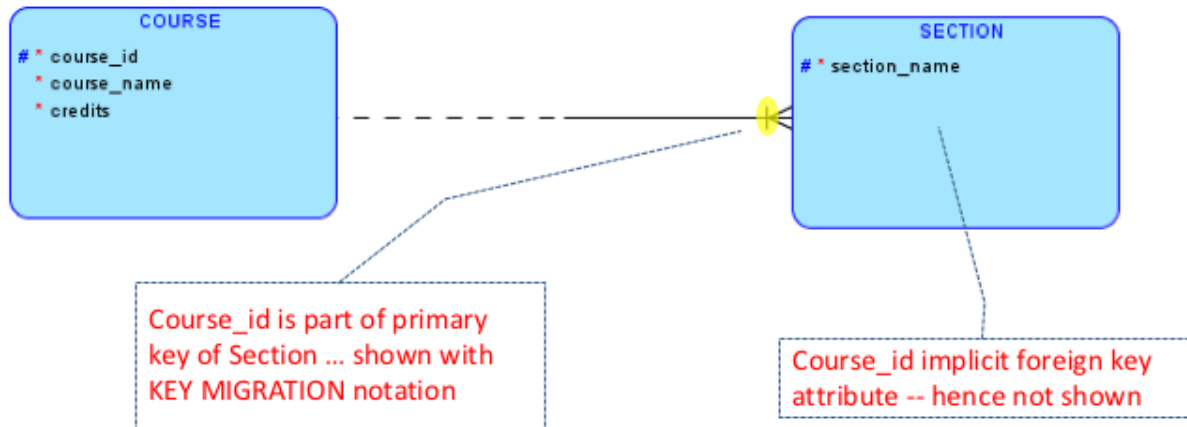
You can easily convince yourself that adding the foreign key to the entity type on the 1 side will not work.

For example, in the above case, adding player\_id as a foreign key to the teams table will not work, because a team could have many players and adding a foreign key will allow us to only show one player per team.

However, adding the team\_id to the players table works because each player can belong to at most one team and adding a foreign key column allows this.

## **ODM Demo**

## Recap 3 – Key Migration



Don't confuse the implicit foreign key attribute with the key migration notation.

The implicit foreign key attribute always arises in all 1:n relationships.

We use the key migration notation only when the foreign key also happens to be part of the primary key in the table on the n side of the relationship. Be sure to understand this from the many examples provided earlier.

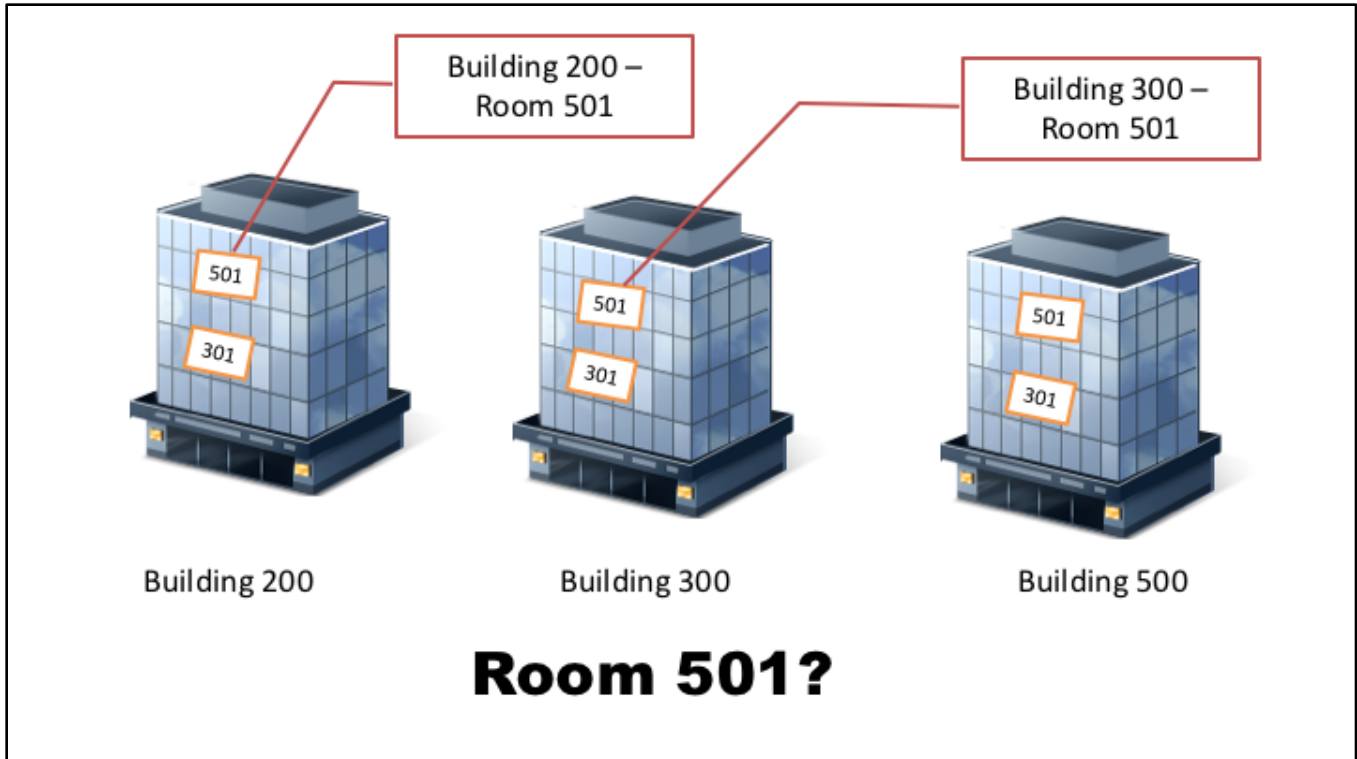
A company has many buildings and each has many rooms. To track the assets installed in each room, the company needs to track all rooms. The company has assigned each building a unique `building_id` and stores the square footage, GPS coordinate and height of each building.

Each building has many rooms and each room has a room number. Room numbers are unique within a building, but not across buildings. Thus building 200 could have a room 103 and building 300 could also have a room 103.


To uniquely identify a room across the company we have to use both the building id and the room number.

A company wants to track the assets installed in various rooms of its office buildings. Hence it needs to have a way to store information about buildings and rooms.






The same room number can occur in several buildings and hence the room number alone does not suffice to uniquely identify a room across all the buildings.



The combination of room number and building number provides a unique identifier for a room. The combination will be unique across all rooms of all buildings.




Key migration notation

Primary key is room\_number + building\_id

## **Key Migration Demo in ODM**


## Relationship Naming



**Two relationships? What do they stand for?**

In situations where we know the application domain very well, we might not need to be told what each relationship represents. However, in unfamiliar domains and sometimes even in familiar domains, we might not understand what a relationship signifies. In the above ERD, we might be a little confused by the two relationships between Player and Team.

To clarify the meaning of relationships, we need to name them.



**Now we see why we have two relationships**  
**Need names to clarify ERDs**


The names on the relationships tell us what each relationship signifies from the viewpoint of the application domain.

One shows that a player is a regular player on a team and the other tells us that each team has one player as captain.

We place names on relationship lines to indicate the role that each entity type plays in the relationship.


## Reading Relationships

Reading from left ...




## Reading Relationships

Reading from right ...





## Reading Relationships



Reading from left ...


**Each PLAYER MIGHT BE ON ONE TEAM**

Reading from right ...

**Each TEAM MIGHT BE COMPRISED OF ONE OR MORE PLAYERS**


## Naming Relationships

Reading from left to right – name the relationship on the left side so as to properly complete the following sentence:



Solid or dashed on left

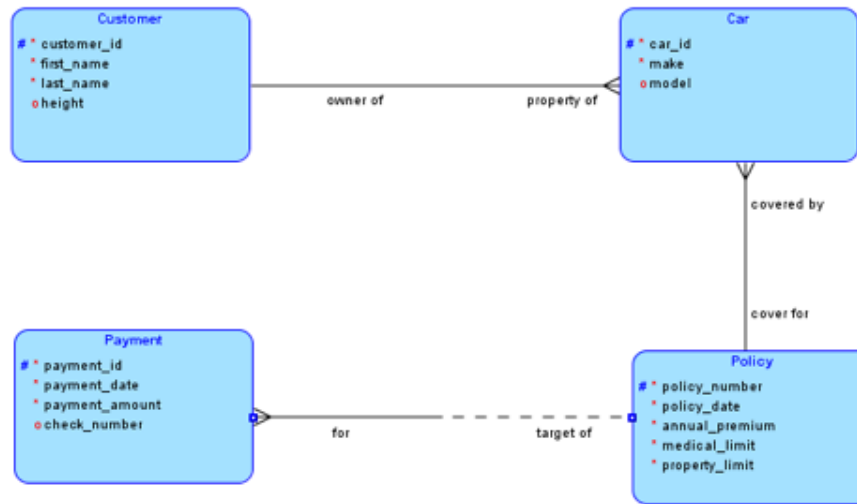
Crow-foot or not on right



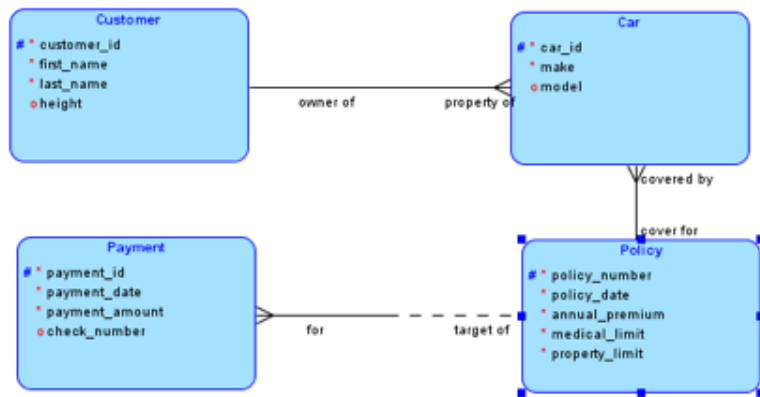
Name-1



## Reading Relationships – Your turn



## Reading Relationships – Your turn



Each customer must be owner of one or many cars

Each car must be property of one customer


Each car must be covered by one policy

Each policy must be the cover for one or more cars


Each policy might be the target of one or more payments

Each payment must be for one policy

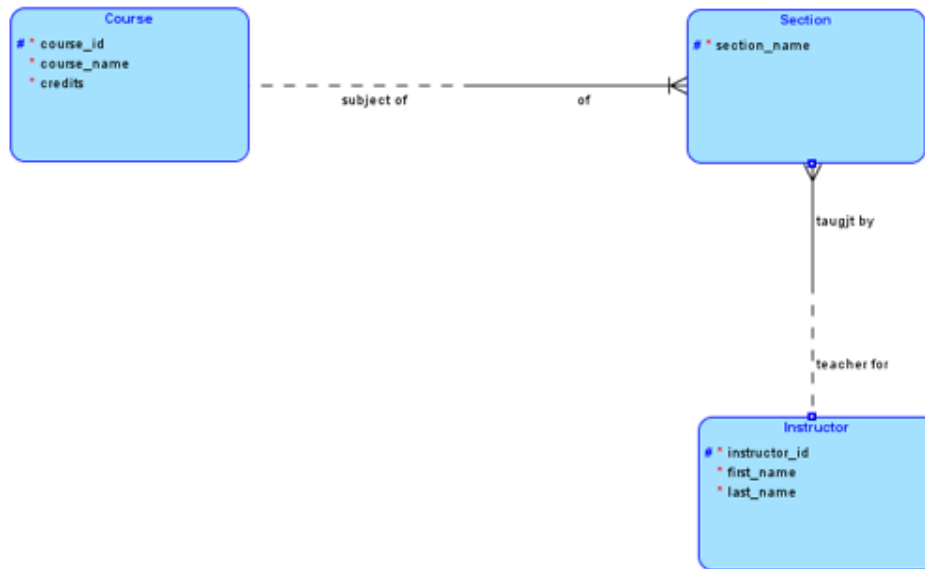
**Naming Relationships – Your Turn: Find suitable names for the relationship – you need two names, one on each side.**



**Naming Relationships – Your Turn (continued): Find suitable names for the relationships – you need two names per relationship.**



## Naming Relationships – Your Turn (continued): Find suitable names for the relationships – you need two names per relationship.



**Naming Relationships – Your Turn (continued): Find suitable names for the relationships – you need two names per relationship.**



## **Relationship Naming Demo in ODM**

# **ERD Ninjahood**




## **Master m:n Relationships**

**Inventory management scenario:**

- A company keeps many products in stock.
- The company has many warehouses and stores many products in each warehouse.
- Each product could also be stored in several warehouses.
- At various points in time the company could run out of stock of some products.
- At various points in time, some warehouses could be empty.

### Inventory management scenario





Attribute quantity?


Attribute values

# Atomic

Entity instances must have **indivisible** values for each attribute

Non-atomic values in column  
– not allowed in a relational  
design


team_id	team_name	player_id
10	Bulls	10, 20, 40, 70
20	Jackals	30, 90, 200
30	Patriots	110, 50, 60, 130, 620



Where does attribute quantity belong? Product or Warehouse?







Because a product can be in many warehouses, placing *quantity* in Product makes it non-atomic

Because a warehouse can have stocks of many products, placing *quantity* in Warehouse makes it non-atomic

Disk drive is in W1 and W2


product_id	product_name	quantity
10	Disk drive	w1:200, w2:100
...	...	...

W2 has both Disk drives and Webcams

warehouse_id	warehouse_name	products
w2	Alpine	10:100, 30:400
...	...	...




Where does attribute  
quantity belong? Product or  
Warehouse?



**Quantity** is an attribute of the relationship between Product and Warehouse

## Duck test



If it swims like a duck and  
quacks like a duck it must  
be a duck

*If a relationship has its own attribute, then it must be an entity type*


New entity type that has ***quantity*** as an attribute

**ALWAYS** create new entity type for m:n relationship


**ALWAYS create new  
entity type for m:n  
relationship ...**

**Even if it does not  
seem to have an  
attribute.**


## Inventory management scenario









Each instance represents the storage of a certain quantity of a product at a warehouse – ***Stock***






**Stock instance:** of a particular product at a particular warehouse

**product\_id + warehouse\_id** can be primary key for Stock



Cardinality notation?

## Relationship Degree




At most how many stocks can a product have?

**n**

At most how many products can a stock be related to?

**1**

**1:n**



At **least** how many stocks must a product have?

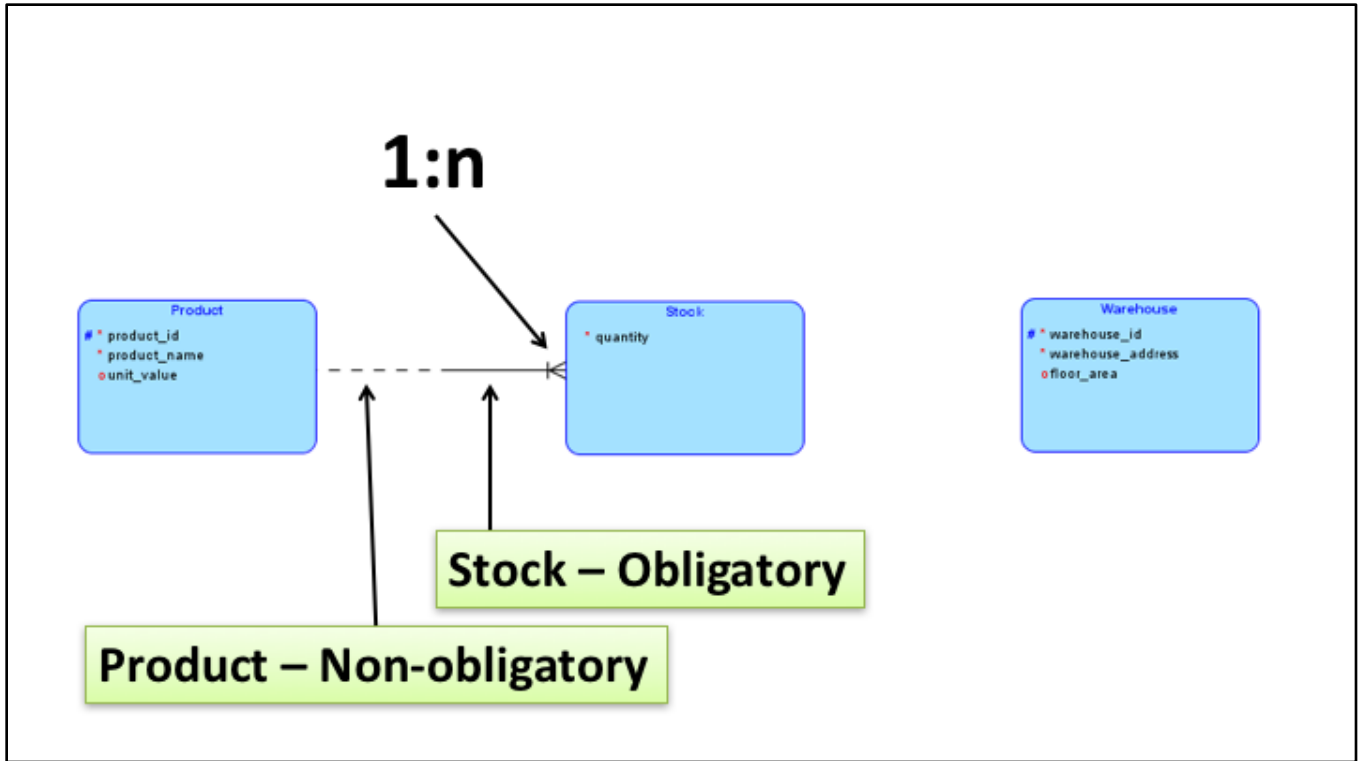
**0**

At **least** how many products must a sock be related to?

**1**

**Product – Non-obligatory**

**Stock – Obligatory**



## Relationship Degree



At most how many stocks can a warehouse have?

**n**

At most how many warehouses can a stock be related to?

**1**

**1:n**





At **least** how many stocks must a warehouse have?

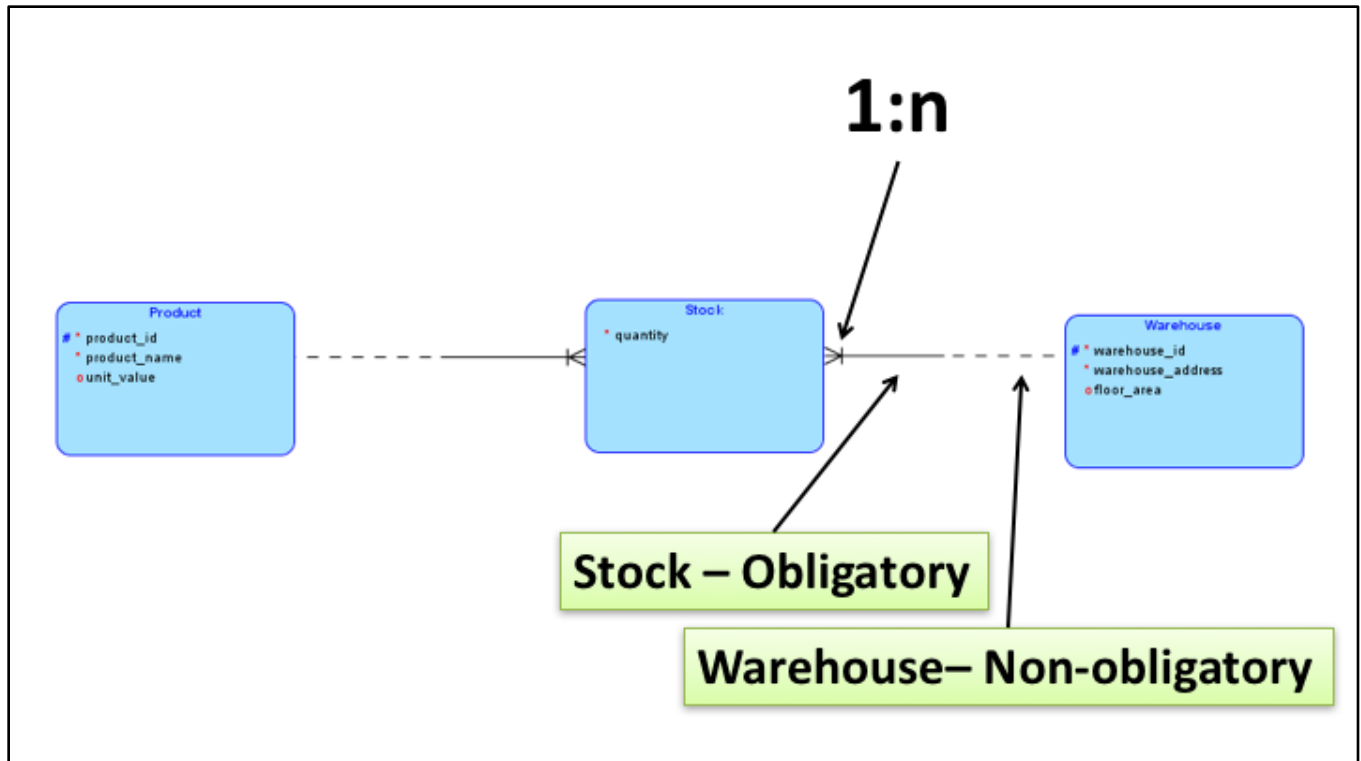
**0**

At **least** how many warehouses must a stock be related to?


**1**

**Warehouse– Non-obligatory**

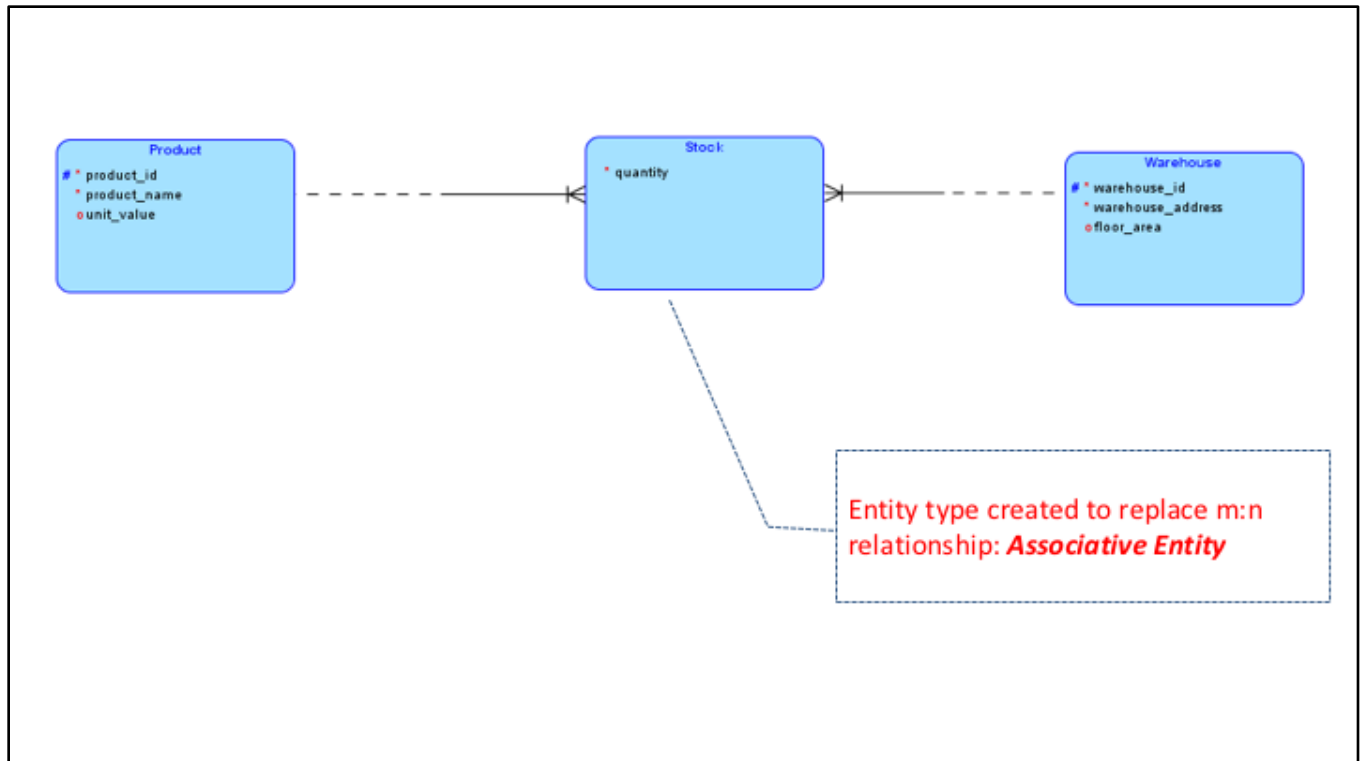
**Stock – Obligatory**




## Primary Key for Stock?



From key migrations: `product_id + warehouse_id`





Cardinality notation on associative entities always the same: crow-feet and solid half-lines on both sides.

Key migrations from both ends – if we do not give it its own primary key

### Hollywood scenario:

- A film studio produces many films.
- The studio has signed on many actors.
- Each actor might have a role on several films.
- Each film must involve one or more actors playing roles.
- Some actors might not yet have acted in any films.



